

Relational Database Design

Charles Severance

www.pg4e.com

<http://www.pg4e.com/lectures/02-Database-Design-Many-to-Many.txt>

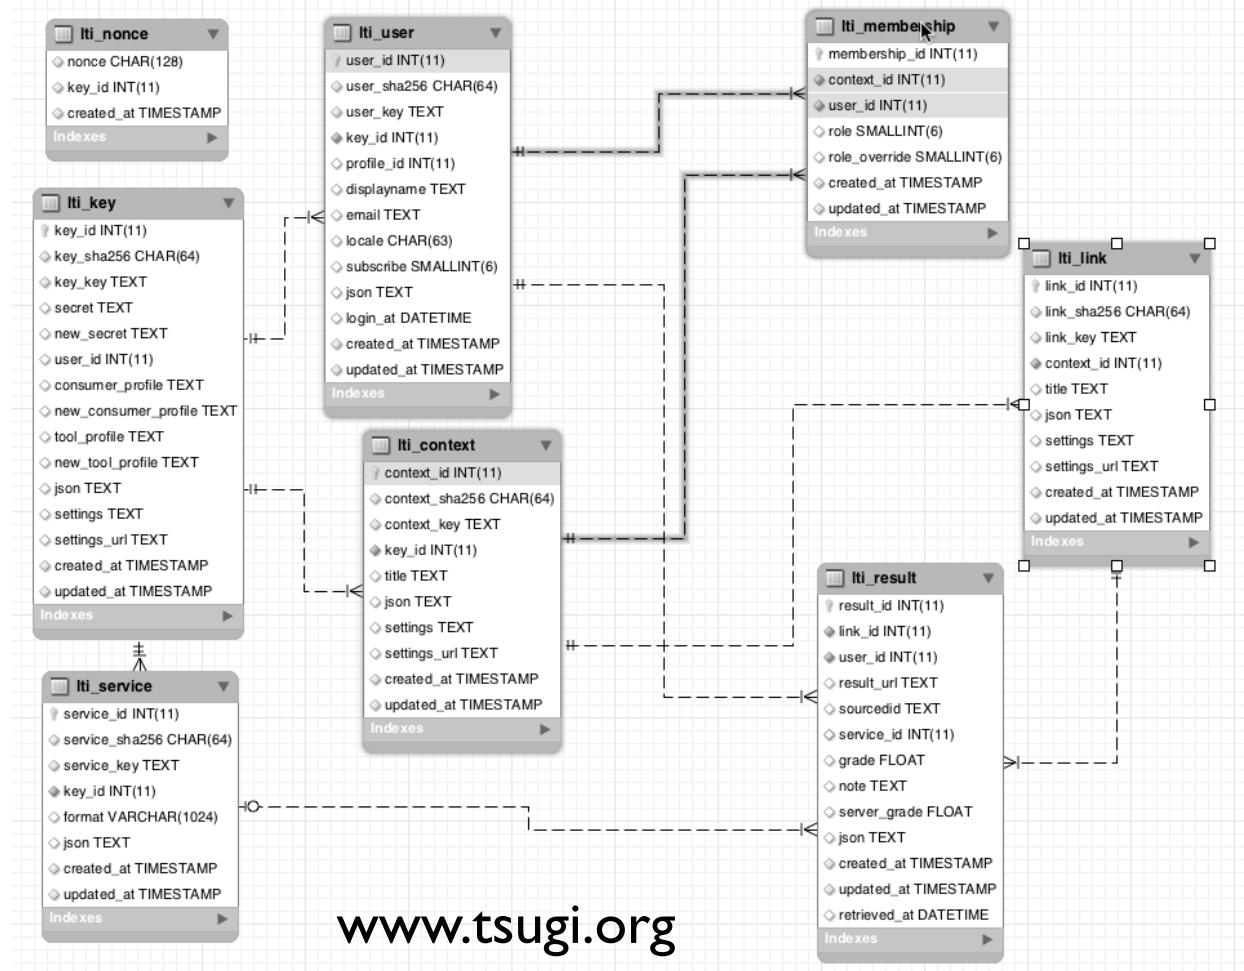


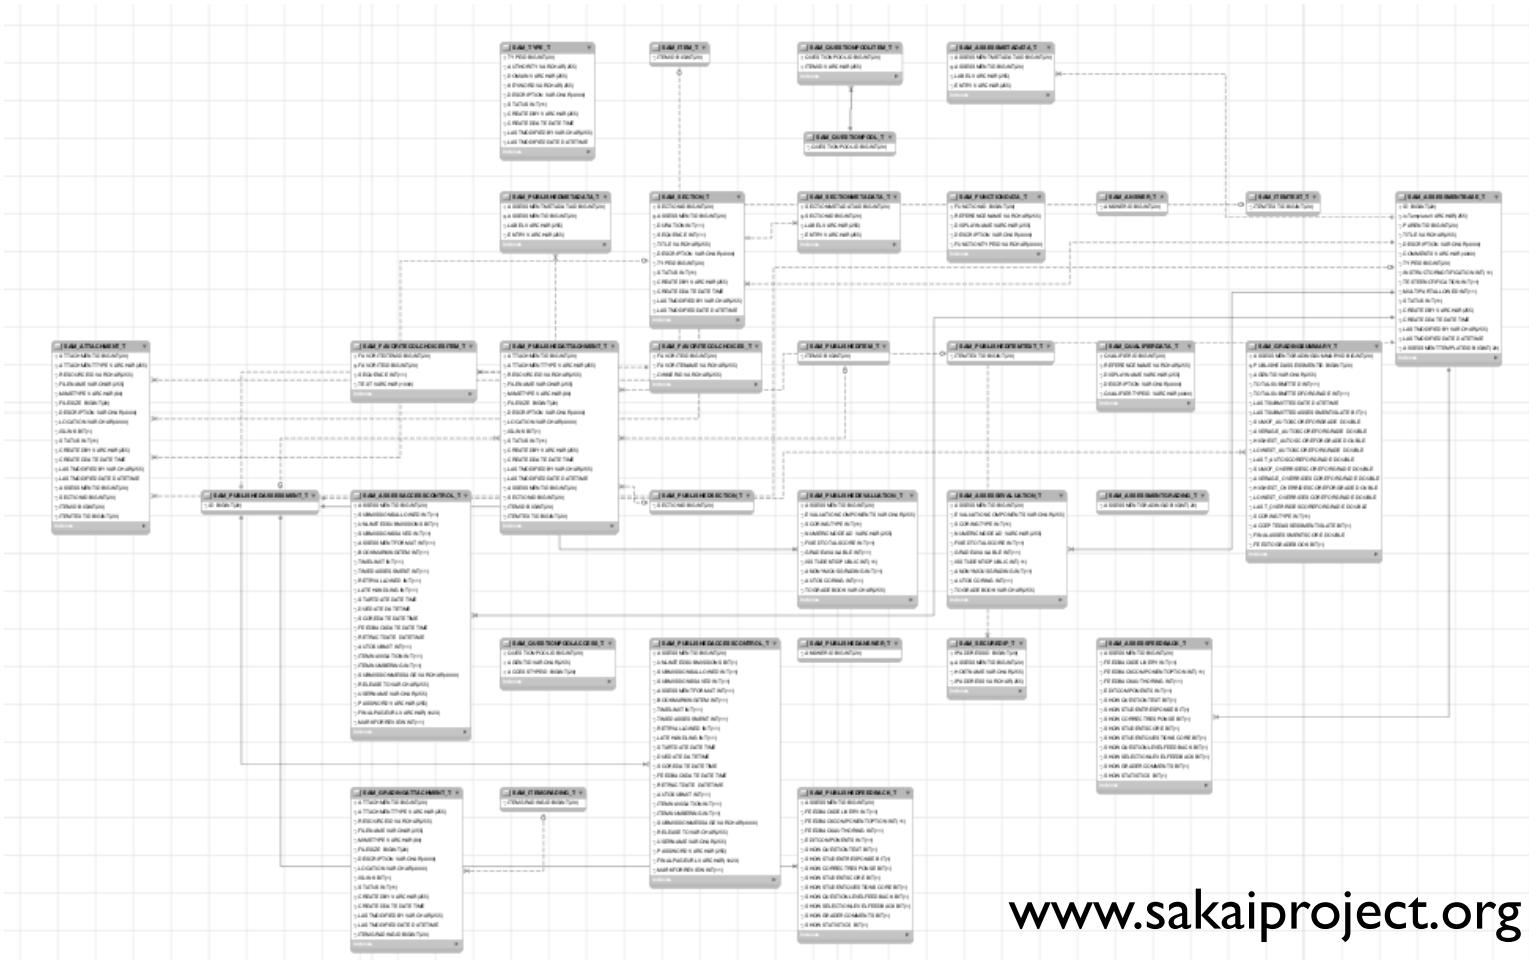
Relational Database Design

http://en.wikipedia.org/wiki/Relational_model

Database Design

- Database design is an art form of its own with particular skills and experience.
- Our goal is to avoid the really bad mistakes and design clean and easily understood databases.
- Others may performance tune things later.
- Database design starts with a picture...





www.sakaiproject.org

Building a Data Model

- Drawing a picture of the data objects for our application and then figuring out how to represent the objects and their relationships
- Basic Rule: Don't put the same string data in twice - use a relationship instead
- When there is one thing in the “real world” there should only be one copy of that thing in the database

Track	Len	Artist	Album	Genre	Rating	Count
<input checked="" type="checkbox"/> Hells Bells	5:13	AC/DC	Who Made Who	Rock	★★★★★	61
<input checked="" type="checkbox"/> Shake Your Foundations	3:54	AC/DC	Who Made Who	Rock	★★★★★	70
<input checked="" type="checkbox"/> Chase the Ace	3:01	AC/DC	Who Made Who	Rock		56
<input checked="" type="checkbox"/> For Those About To Rock (We ...	5:54	AC/DC	Who Made Who	Rock	★★★★★	61
<input checked="" type="checkbox"/> Dúlamán	3:43	Altan	Natural Wonders M...	New Age		31
<input checked="" type="checkbox"/> Rode Across the Desert	4:10	America	Greatest Hits	Easy Listen...	★★★★★	23
<input checked="" type="checkbox"/> Now You Are Gone	3:08	America	Greatest Hits	Easy Listen...	★★★★★	18
<input checked="" type="checkbox"/> Tin Man	3:30	America	Greatest Hits	Easy Listen...	★★★★★	23
<input checked="" type="checkbox"/> Sister Golden Hair	3:22	America	Greatest Hits	Easy Listen...	★★★★★	24
<input checked="" type="checkbox"/> Track 01	4:22	Billy Price	Danger Zone	Blues/R&B	★★★★★	26
<input checked="" type="checkbox"/> Track 02	2:45	Billy Price	Danger Zone	Blues/R&B	★★★★★	18
<input checked="" type="checkbox"/> Track 03	3:26	Billy Price	Danger Zone	Blues/R&B	★★★★★	22
<input checked="" type="checkbox"/> Track 04	4:17	Billy Price	Danger Zone	Blues/R&B	★★★★★	18
<input checked="" type="checkbox"/> Track 05	3:50	Billy Price	Danger Zone	Blues/R&B	★★★★★	21
<input checked="" type="checkbox"/> War Pigs/Luke's Wall	7:58	Black Sabbath	Paranoid	Metal	★★★★★	25
<input checked="" type="checkbox"/> Paranoid	2:53	Black Sabbath	Paranoid	Metal	★★★★★	22
<input checked="" type="checkbox"/> Planet Caravan	4:35	Black Sabbath	Paranoid	Metal	★★★★★	25
<input checked="" type="checkbox"/> Iron Man	5:59	Black Sabbath	Paranoid	Metal	★★★★★	26
<input checked="" type="checkbox"/> Electric Funeral	4:53	Black Sabbath	Paranoid	Metal	★★★★★	22
<input checked="" type="checkbox"/> Hand of Doom	7:10	Black Sabbath	Paranoid	Metal	★★★★★	23
<input checked="" type="checkbox"/> Rat Salad	2:30	Black Sabbath	Paranoid	Metal	★★★★★	31
<input checked="" type="checkbox"/> Jack the Stripper/Fairies Wear ...	6:14	Black Sabbath	Paranoid	Metal	★★★★★	24
<input checked="" type="checkbox"/> Bomb Squad (TECH)	3:28	Brent	Brent's Album			1
<input checked="" type="checkbox"/> clay techno	4:36	Brent	Brent's Album			2
<input checked="" type="checkbox"/> Heavy	3:08	Brent	Brent's Album			1
<input checked="" type="checkbox"/> Hi metal man	4:20	Brent	Brent's Album			1
<input checked="" type="checkbox"/> Mistro	2:58	Brent	Brent's Album			1

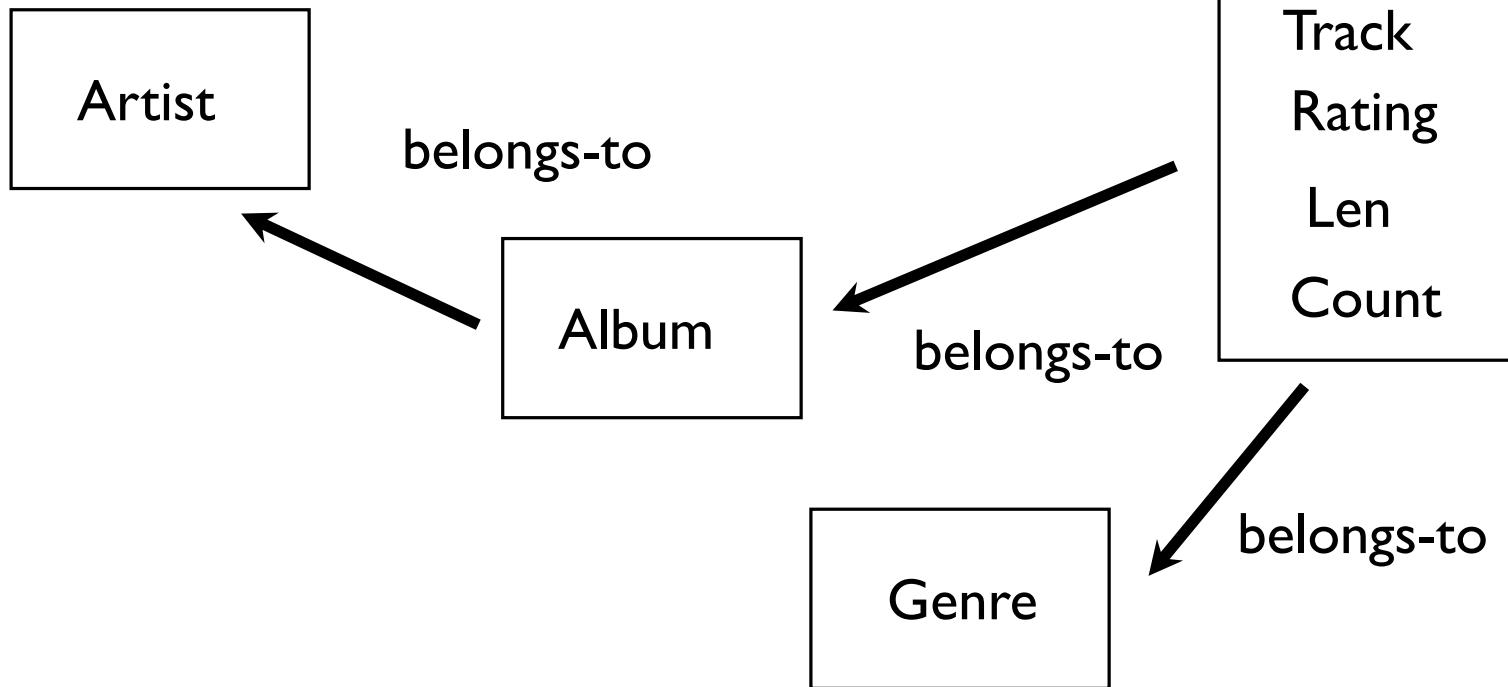
For each “piece of info”...

- Is the column an object or an attribute of another object?
- Once we define objects, we need to define the relationships between objects.

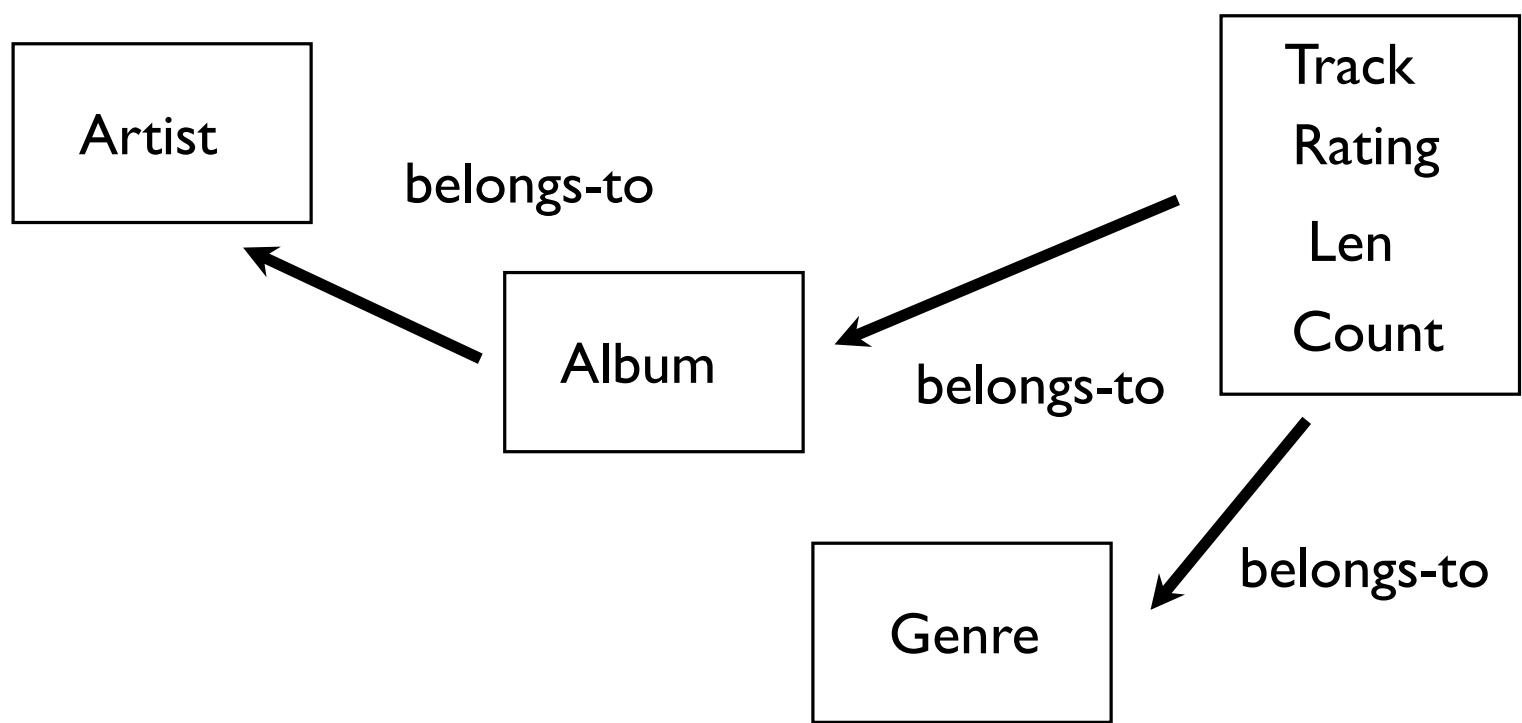
Len	Album
Genre	Artist
Rating	Track
Count	

✓ Hells Bells	5:13	AC/DC	Who Made Who	Rock	★★★★★	61
✓ Shake Your Foundations	3:54	AC/DC	Who Made Who	Rock	★★★★★	70
✓ Chase the Ace	3:01	AC/DC	Who Made Who	Rock	★★★★★	56
✓ For Those About To Rock (We ...	5:54	AC/DC	Who Made Who	Rock	★★★★★	61
✓ Dúlamán	3:43	Altan	Natural Wonders M...	New Age	★★★★★	31
✓ Rode Across the Desert	4:10	America	Greatest Hits	Easy Listen...	★★★★★	23
✓ Now You Are Gone	3:08	America	Greatest Hits	Easy Listen...	★★★★★	18
✓ Tie Me	3:20	America	Greatest Hits	Easy Listen...	★★★★★	22

Track
Album
Artist
Genre
Rating
Len
Count



<input checked="" type="checkbox"/> Hells Bells	5:13	AC/DC	Who Made Who	Rock	★★★★★	61
<input checked="" type="checkbox"/> Shake Your Foundations	3:54	AC/DC	Who Made Who	Rock	★★★★★	70
<input checked="" type="checkbox"/> Chase the Ace	3:01	AC/DC	Who Made Who	Rock	★★★★★	56
<input checked="" type="checkbox"/> For Those About To Rock (We ...	5:54	AC/DC	Who Made Who	Rock	★★★★★	61
<input checked="" type="checkbox"/> Dúlamán	3:43	Altan	Natural Wonders M...	New Age	★★★★★	31
<input checked="" type="checkbox"/> Rode Across the Desert	4:10	America	Greatest Hits	Easy Listen...	★★★★★	23
<input checked="" type="checkbox"/> Now You Are Gone	3:08	America	Greatest Hits	Easy Listen...	★★★★★	18
<input checked="" type="checkbox"/> The Man	3:20	America	Greatest Hits	Easy Listen...	★★★★★	22



<input checked="" type="checkbox"/> Hells Bells	5:13	AC/DC	Who Made Who	Rock	★★★★★	61
<input checked="" type="checkbox"/> Shake Your Foundations	3:54	AC/DC	Who Made Who	Rock	★★★★★	70
<input checked="" type="checkbox"/> Chase the Ace	3:01	AC/DC	Who Made Who	Rock	★★★★★	56
<input checked="" type="checkbox"/> For Those About To Rock (We ...	5:54	AC/DC	Who Made Who	Rock	★★★★★	61
<input checked="" type="checkbox"/> Dúlamán	3:43	Altan	Natural Wonders M...	New Age	★★★★★	31
<input checked="" type="checkbox"/> Rode Across the Desert	4:10	America	Greatest Hits	Easy Listen...	★★★★★	23
<input checked="" type="checkbox"/> Now You Are Gone	3:08	America	Greatest Hits	Easy Listen...	★★★★★	18
<input checked="" type="checkbox"/> The Man	3:20	America	Greatest Hits	Easy Listen...	★★★★★	22

Key Terminology

Finding our way around....

Three Kinds of Keys

- Primary key - generally an integer auto-increment field
- Logical key - what the outside world uses for lookup
- Foreign key - generally an integer key pointing to a row in another table

Album
album_id
title
artist_id
...

Primary Key Rules

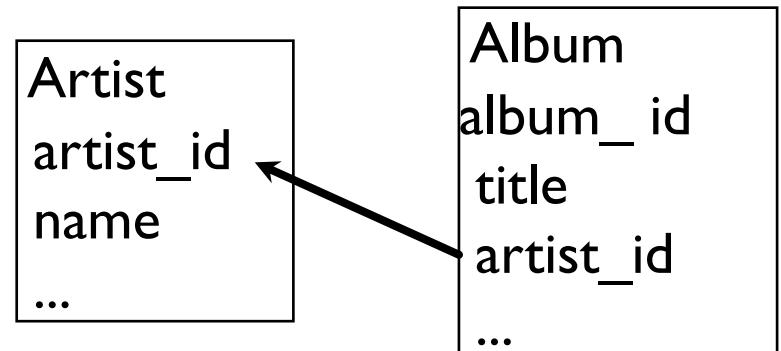
Best practices:

- Never use your logical key as the primary key.
- Logical keys can and do change, albeit slowly.
- Relationships that are based on matching string fields are less efficient than integers.

User
user_id
email
password
name
created_at
modified_at
login_at

Foreign Keys

- A foreign key is when a table has a column containing a key that points to the primary key of another table.
- When all primary keys are integers, then all foreign keys are integers. This is good - very good.



Normalization and Foreign Keys

<input checked="" type="checkbox"/> Hells Bells	5:13	AC/DC	Who Made Who	Rock		61
<input checked="" type="checkbox"/> Shake Your Foundations	3:54	AC/DC	Who Made Who	Rock		70
<input checked="" type="checkbox"/> Chase the Ace	3:01	AC/DC	Who Made Who	Rock		56
<input checked="" type="checkbox"/> For Those About To Rock (We ...	5:54	AC/DC	Who Made Who	Rock		61
<input checked="" type="checkbox"/> Dúlamán	3:43	Altan	Natural Wonders M...	New Age		31
<input checked="" type="checkbox"/> Rode Across the Desert	4:10	America	Greatest Hits	Easy Listen...		23
<input checked="" type="checkbox"/> Now You Are Gone	3:08	America	Greatest Hits	Easy Listen...		18
<input type="checkbox"/> Tin Man	2:20	America	Contact Hits	Easy Listen...		22

We want to keep track of which band is the “creator” of each music track...
 What album does this song “belong to”?

Which album is this song related to?

Database Normalization (3NF)

There is *tons* of database theory - way too much to understand without excessive predicate calculus

- Do not replicate data. Instead, reference data. Point at data.
- Use integers for keys and for references.
- Add a special “key” column to each table, which you will make references to.

http://en.wikipedia.org/wiki/Database_normalization

Integer Reference Pattern

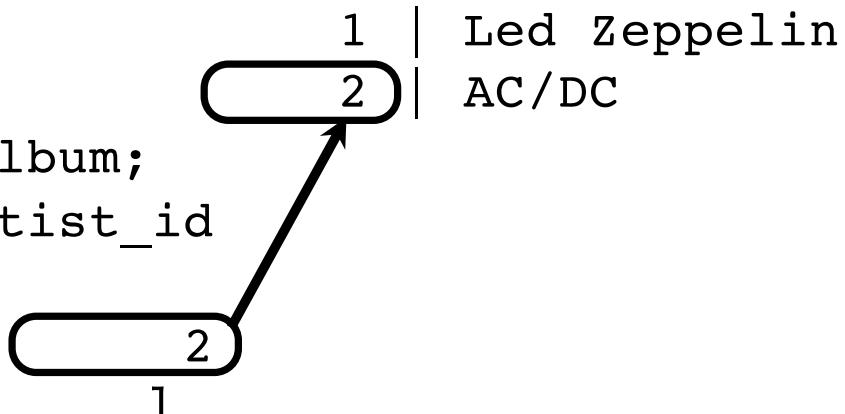
We use integer columns in one table to reference (or look up) rows in another table.

```
music=> SELECT * FROM album;
```

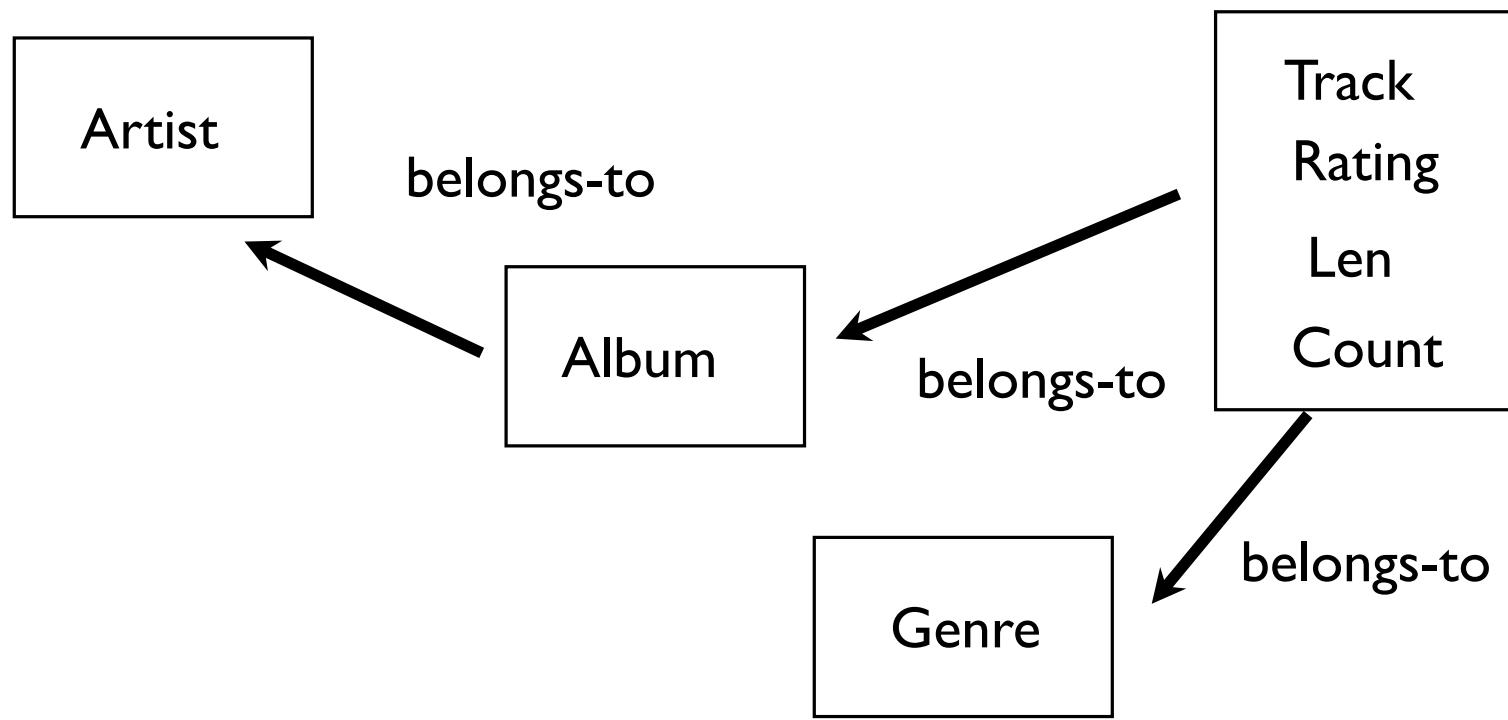
id	title	artist_id	
-----+-----+-----+			
1	Who Made Who		
2	IV		

```
music=> SELECT * FROM artist;
```

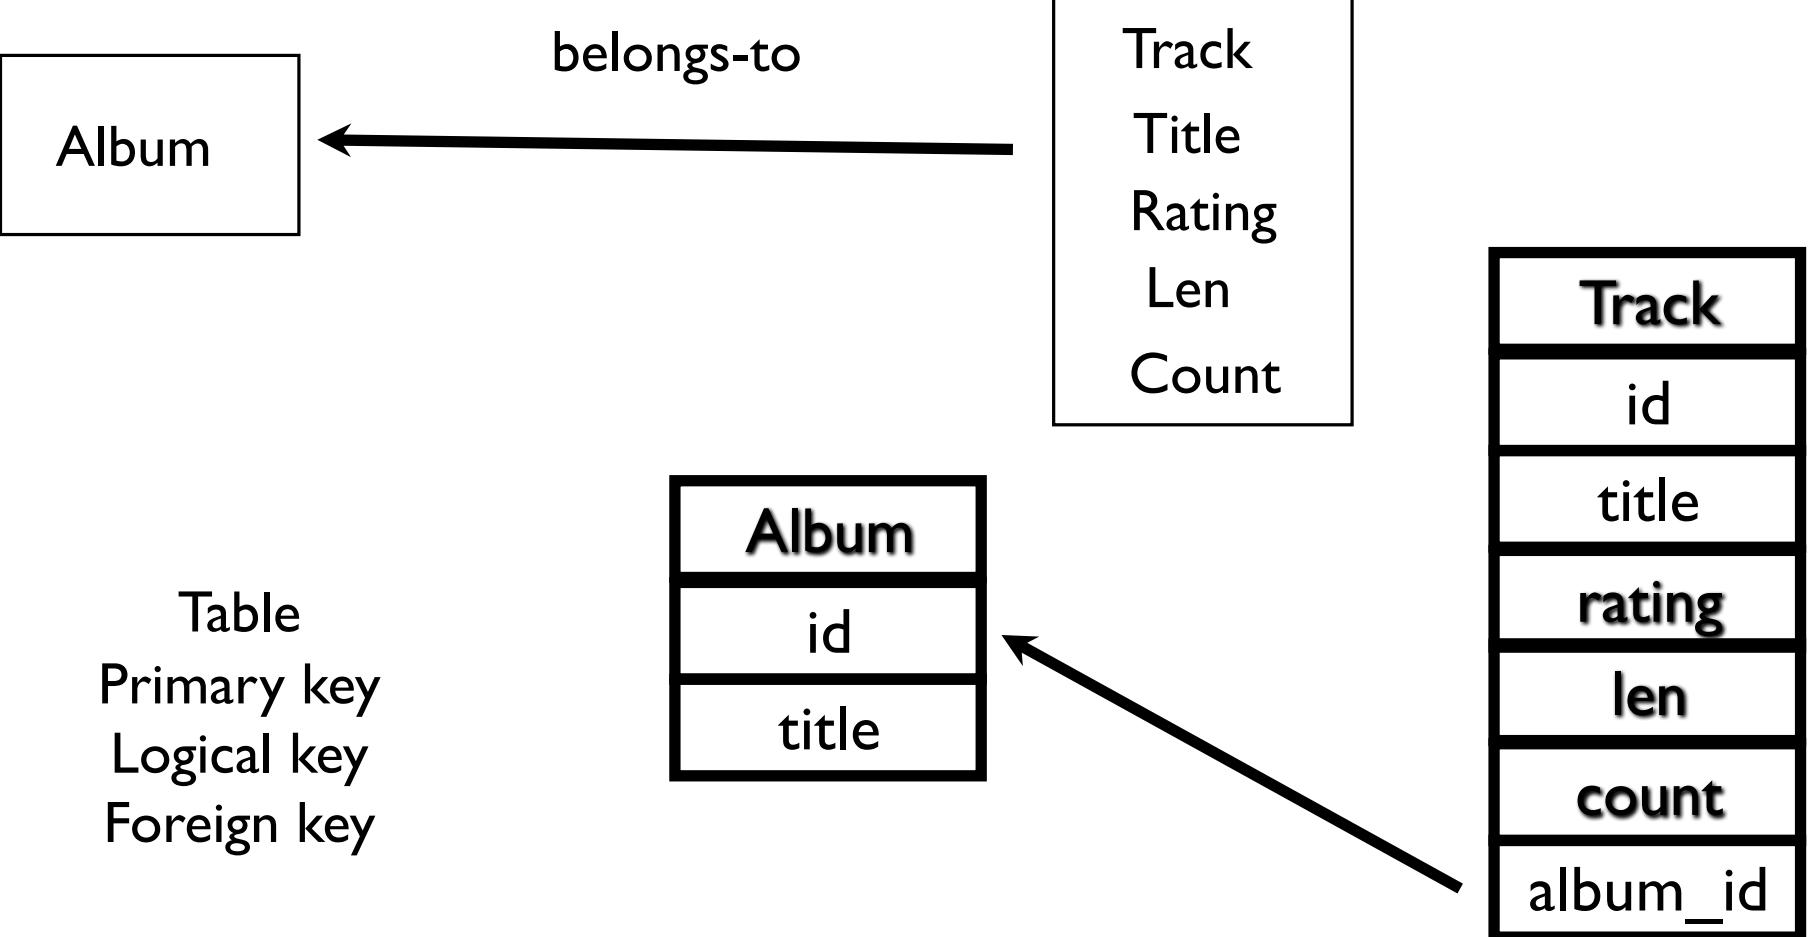
id	name
-----+-----	
1	Led Zeppelin
2	AC/DC

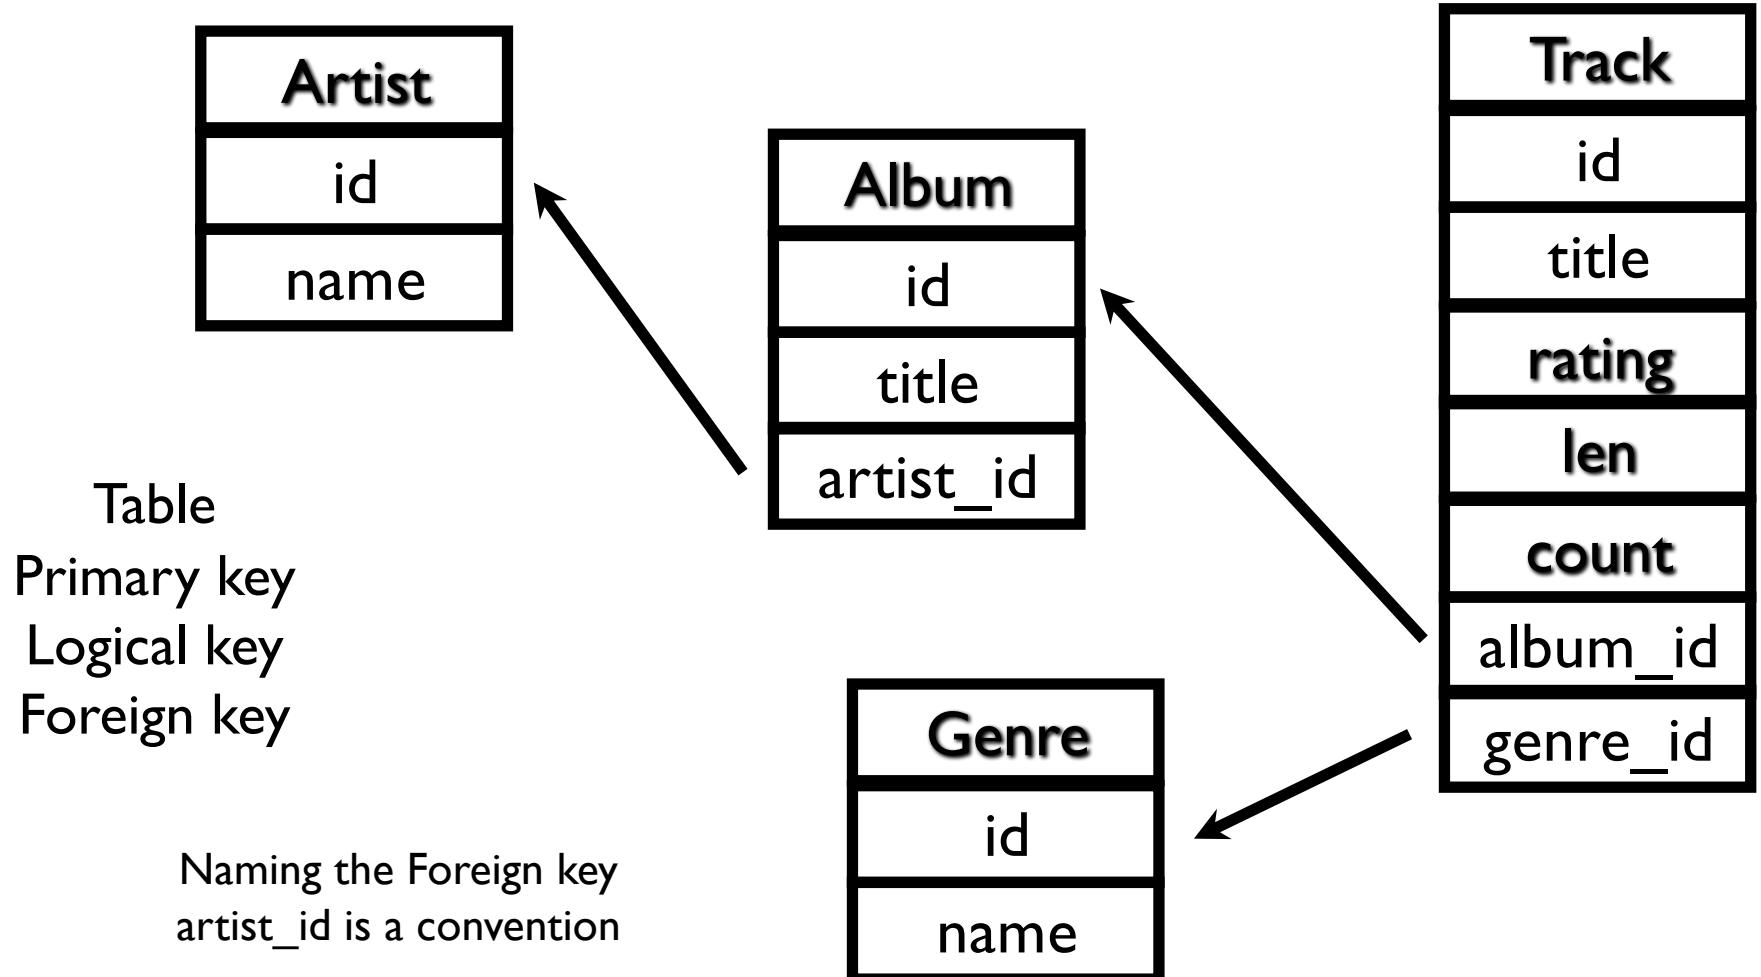


Building a Physical Data Schema



<input checked="" type="checkbox"/> Hells Bells	5:13	AC/DC	Who Made Who	Rock	★★★★★	61
<input checked="" type="checkbox"/> Shake Your Foundations	3:54	AC/DC	Who Made Who	Rock	★★★★★	70
<input checked="" type="checkbox"/> Chase the Ace	3:01	AC/DC	Who Made Who	Rock		56
<input checked="" type="checkbox"/> For Those About To Rock (We ...	5:54	AC/DC	Who Made Who	Rock	★★★★★	61
<input checked="" type="checkbox"/> Dúlamán	3:43	Altan	Natural Wonders M...	New Age		31
<input checked="" type="checkbox"/> Rode Across the Desert	4:10	America	Greatest Hits	Easy Listen...	★★★★★	23
<input checked="" type="checkbox"/> Now You Are Gone	3:08	America	Greatest Hits	Easy Listen...	★★★★★	18
<input checked="" type="checkbox"/> Tie Me	3:20	America	Greatest Hits	Easy Listen...	★★★★★	22





Creating our Music Database

```
sudo -u postgres psql postgres
```

```
postgres=# CREATE DATABASE music
          WITH OWNER 'pg4e' ENCODING 'UTF8';
CREATE DATABASE
postgres=#
```

```
CREATE TABLE artist (
    id SERIAL,
    name VARCHAR(128) UNIQUE,
    PRIMARY_KEY(id)
);

CREATE TABLE album (
    id SERIAL,
    title VARCHAR(128) UNIQUE,
    artist_id INTEGER REFERENCES artist(id) ON DELETE CASCADE,
    PRIMARY KEY(id)
);
```

```
CREATE TABLE genre (
    id SERIAL,
    name VARCHAR(128) UNIQUE,
    PRIMARY_KEY(id)
);

CREATE TABLE track (
    id SERIAL,
    title VARCHAR(128),
    len INTEGER,
    rating INTEGER,
    count INTEGER,
    album_id INTEGER REFERENCES genre(id) ON DELETE CASCADE,
    genre_id INTEGER REFERENCES album(id) ON DELETE CASCADE,
    UNIQUE(title, album_id),
    PRIMARY KEY(id)
);
```

```
music=> \d track
                                         Table "public.track"
  Column |          Type          |                         Modifiers
-----+-----+-----+
    id   |      integer      | not null default nextval('track_id_seq'::regclass)
  title | character varying(128) |
    len  |      integer      |
  rating |      integer      |
  count |      integer      |
album_id |      integer      |
genre_id |      integer      |
Indexes:
  "track_pkey" PRIMARY KEY, btree (id)
  "track_title_album_id_key" UNIQUE CONSTRAINT, btree (title, album_id)
Foreign-key constraints:
  "track_album_id_fkey" FOREIGN KEY (album_id) REFERENCES album(id) ON DELETE CASCADE
  "track_genre_id_fkey" FOREIGN KEY (genre_id) REFERENCES genre(id) ON DELETE CASCADE
```

```
music=>
```

```
music=> INSERT INTO artist (name) VALUES ('Led Zeppelin');
INSERT 0 1
music=> INSERT INTO artist (name) VALUES ('AC/DC');
INSERT 0 1
music=> SELECT * FROM artist;
+----+-----+
| id | name |
+----+-----+
| 1  | Led Zeppelin |
| 2  | AC/DC |
(2 rows)

music=>
```

```
music=> INSERT INTO album (title, artist_id) VALUES ('Who Made Who', 2);
INSERT 0 1
music=> INSERT INTO album (title, artist_id) VALUES ('IV', 1);
INSERT 0 1
music=> SELECT * FROM album;
+----+-----+-----+
| id | title | artist_id |
+----+-----+-----+
| 1  | Who Made Who |      2
| 2  | IV           |      1
+----+-----+-----+
(2 rows)
```

```
music=> INSERT INTO genre (name) VALUES ('Rock');
INSERT 0 1
music=> INSERT INTO genre (name) VALUES ('Metal');
INSERT 0 1
music=> SELECT * FROM genre;
+----+-----+
| id | name |
+----+-----+
| 1  | Rock |
| 2  | Metal|
+----+
(2 rows)
```

```
music=> INSERT INTO track (title, rating, len, count, album_id, genre_id)
music->      VALUES ('Black Dog', 5, 297, 0, 2, 1) ;
INSERT 0 1
music=> INSERT INTO track (title, rating, len, count, album_id, genre_id)
music->      VALUES ('Stairway', 5, 482, 0, 2, 1) ;
INSERT 0 1
music=> INSERT INTO track (title, rating, len, count, album_id, genre_id)
music->      VALUES ('About to Rock', 5, 313, 0, 1, 2) ;
INSERT 0 1
music=> INSERT INTO track (title, rating, len, count, album_id, genre_id)
music->      VALUES ('Who Made Who', 5, 207, 0, 1, 2) ;
INSERT 0 1
music=> SELECT * FROM track;
+----+-----+-----+-----+-----+-----+
| id | title | len | rating | count | album_id | genre_id |
+----+-----+-----+-----+-----+-----+
| 1  | Black Dog | 297 | 5 | 0 | 2 | 1 |
| 2  | Stairway | 482 | 5 | 0 | 2 | 1 |
| 3  | About to Rock | 313 | 5 | 0 | 1 | 2 |
| 4  | Who Made Who | 207 | 5 | 0 | 1 | 2 |
(4 rows)
```

```
music=> SELECT * FROM track;
+-----+
| id | title | len | rating | count | album_id | genre_id |
+-----+
| 1 | Black Dog | 297 | 5 | 0 | 2 | 1 |
| 2 | Stairway | 482 | 5 | 0 | 2 | 1 |
| 3 | About to Rock | 313 | 5 | 0 | 1 | 2 |
| 4 | Who Made Who | 207 | 5 | 0 | 1 | 2 |
+-----+
```

```
music=> SELECT * FROM genre;
+-----+
| id | name |
+-----+
| 1 | Rock |
| 2 | Metal |
+-----+
```

```
music=> SELECT * FROM album;
+-----+
| id | title | artist_id |
+-----+
| 1 | Who Made Who | 1 |
| 2 | IV | 2 |
+-----+
```

```
music=> SELECT * FROM artist;
+-----+
| id | name |
+-----+
| 1 | Led Zeppelin |
| 2 | AC/DC |
+-----+
```

We Have Relationships!

Using Join Across Tables

[http://en.wikipedia.org/wiki/Join_\(SQL\)](http://en.wikipedia.org/wiki/Join_(SQL))

Relational Power

- By removing the replicated data and replacing it with references to a single copy of each bit of data, we build a “web” of information that the relational database can read through very quickly - even for very large amounts of data.
- Often when you want some data it comes from a number of tables linked by these foreign keys.

The JOIN Operation

- The JOIN operation links across several tables as part of a SELECT operation.
- You must tell the JOIN how to use the keys that make the connection between the tables using an ON clause.

```
music=> SELECT * FROM album;
```

id	title	artist_id
1	Who Made Who	1
2	IV	2

```
music=> SELECT * FROM artist;
```

id	name
1	Led Zeppelin
2	AC/DC

1 2
| |
| |
1 2



```
music=> SELECT album.title, artist.name
music->      FROM album JOIN artist
music->      ON album.artist_id = artist.id;
```

title	name
Who Made Who	AC/DC
IV	Led Zeppelin

What we want to see
The tables that hold the data
How the tables are linked

```
music=> SELECT * FROM album;
```

id	title	artist_id
1	Who Made Who	2
2	IV	

```
music=> SELECT * FROM artist;
```

id	name
1	Led Zeppelin
2	AC/DC



```
music=> SELECT album.title, album.artist_id, artist.id, artist.name  
music->      FROM album INNER JOIN artist ON album.artist_id = artist.id;
```

title	artist_id	id	name
Who Made Who	2	2	AC/DC
IV	1	1	Led Zepplin

```
music=> SELECT track.title, track.genre_id, genre.id, genre.name
music->      FROM track CROSS JOIN genre;
    title      | genre_id | id | name
-----+-----+-----+
 Black Dog      |       1 |   1 | Rock
Stairway        |       1 |   1 | Rock
About to Rock   |       2 |   1 | Rock
Who Made Who    |       2 |   1 | Rock
Black Dog       |       1 |   2 | Metal
Stairway        |       1 |   2 | Metal
About to Rock   |       2 |   2 | Metal
Who Made Who    |       2 |   2 | Metal
```

```
music=> SELECT * FROM track;
+-----+-----+-----+-----+-----+-----+
| id   | title    | len | rating | count | album_id | genre_id |
+-----+-----+-----+-----+-----+-----+
| 1    | Black Dog | 297 | 5      | 0     | 2        | 1         |
| 2    | Stairway   | 482 | 5      | 0     | 2        | 1         |
| 3    | About to Rock | 313 | 5      | 0     | 1        | 2         |
| 4    | Who Made Who | 207 | 5      | 0     | 1        | 2         |

```

```
music=> SELECT * FROM genre;
+-----+-----+
| id   | name   |
+-----+-----+
| 1    | Rock   |
| 2    | Metal  |

```

```
music=> SELECT track.title, genre.name
music->      FROM track JOIN genre
music->      ON track.genre_id = genre.id;
+-----+-----+
| title    | name   |
+-----+-----+
| Black Dog | Rock   |
| Stairway   | Rock   |
| About to Rock | Metal |
| Who Made Who | Metal |

```

It Can Get Complex...

```
music=> SELECT track.title, artist.name, album.title, genre.name  
music-> FROM track  
music->     JOIN genre ON track.genre_id = genre.id  
music->     JOIN album ON track.album_id = album.id  
music->     JOIN artist ON album.artist_id = artist.id;
```

title	name	title	genre
Black Dog	Led Zeppelin	IV	Rock
Stairway	Led Zeppelin	IV	Rock
About to Rock	AC/DC	Who Made Who	Metal
Who Made Who	AC/DC	Who Made Who	Metal

<input checked="" type="checkbox"/> Hells Bells	5:13	AC/DC	Who Made Who	Rock	★★★★★	61
<input checked="" type="checkbox"/> Shake Your Foundations	3:54	AC/DC	Who Made Who	Rock	★★★★★	70
<input checked="" type="checkbox"/> Chase the Ace	3:01	AC/DC	Who Made Who	Rock		56
<input checked="" type="checkbox"/> For Those About To Rock (We ...	5:54	AC/DC	Who Made Who	Rock	★★★★★	61
<input checked="" type="checkbox"/> Dúlamán	3:43	Altan	Natural Wonders M...	New Age		31
<input checked="" type="checkbox"/> Rode Across the Desert	4:10	America	Greatest Hits	Easy Listen...	★★★★★	23
<input checked="" type="checkbox"/> Now You Are Gone	3:08	America	Greatest Hits	Easy Listen...	★★★★★	18
<input checked="" type="checkbox"/> Tin Man	3:30	America	Greatest Hits	Easy Listen...	★★★★★	23
<input checked="" type="checkbox"/> Sister Golden Hair	3:22	America	Greatest Hits	Easy Listen...	★★★★★	24
<input checked="" type="checkbox"/> Track 01	4:22	Billy Price	Danger Zone	Blues/R&B	★★★★★	26
<input checked="" type="checkbox"/> Track 02	2:45	Billy Price	Danger Zone	Blues/R&B	★★★★★	18
<input checked="" type="checkbox"/> Track 03	3:26	Billy Price	Danger Zone	Blues/R&B	★★★★★	22
<input checked="" type="checkbox"/> Track 04	4:17	Billy Price	Danger Zone	Blues/R&B	★★★★★	18
<input checked="" type="checkbox"/> Track 05	3:50	Billy Price	Danger Zone	Blues/R&B	★★★★★	21

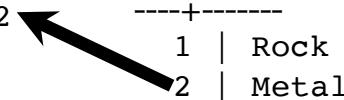
title	name	title	name
Black Dog	Led Zeppelin	IV	Rock
Stairway	Led Zeppelin	IV	Rock
About to Rock	AC/DC	Who Made Who	Metal
Who Made Who	AC/DC	Who Made Who	Metal
4:20	Brent	Brent's Album	1
2:58	Brent	Brent's Album	1

ON DELETE CASCADE

Child

```
music=> SELECT * FROM track;
+-----+
| id | title      | len | rating | count | album_id | genre_id |
+-----+
| 1  | Black Dog   | 297 | 5       | 0      | 2         | 1          |
| 2  | Stairway    | 482 | 5       | 0      | 2         | 1          |
| 3  | About to Rock | 313 | 5       | 0      | 1         | 2          |
| 4  | Who Made Who | 207 | 5       | 0      | 1         | 2          |
```

```
music=> SELECT * FROM genre;
+-----+
| id | name      |
+-----+
| 1  | Rock      |
| 2  | Metal     |
+-----+
```



We are telling Postgres to
"clean up" broken references

Parent

```
DELETE FROM Genre WHERE name = 'Metal'
```

ON DELETE CASCADE

```
music=> SELECT * FROM track;
+-----+-----+-----+-----+-----+-----+-----+
| id | title | len | rating | count | album_id | genre_id |
+-----+-----+-----+-----+-----+-----+-----+
| 1 | Black Dog | 297 | 5 | 0 | 2 | 1 |
| 2 | Stairway | 482 | 5 | 0 | 2 | 1 |
| 3 | About to Rock | 313 | 5 | 0 | 1 | 2 |
| 4 | Who Made Who | 207 | 5 | 0 | 1 | 2 |
+-----+-----+-----+-----+-----+-----+-----+
(4 rows)
```

```
music=> DELETE FROM genre WHERE name='Metal';
```

```
DELETE 1
```

```
music=> SELECT * FROM track;
```

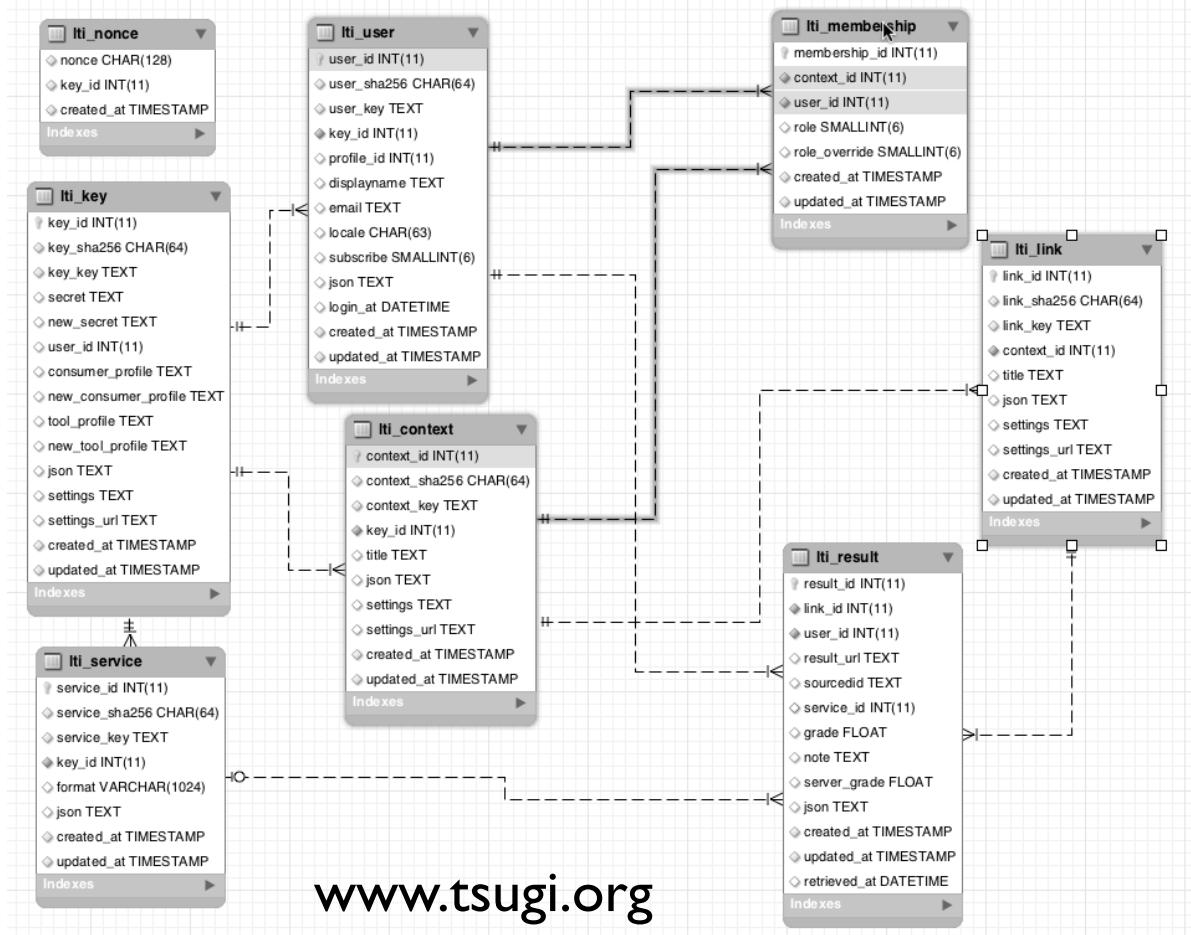
```
+-----+-----+-----+-----+-----+-----+-----+
| id | title | len | rating | count | album_id | genre_id |
+-----+-----+-----+-----+-----+-----+-----+
| 1 | Black Dog | 297 | 5 | 0 | 2 | 1 |
| 2 | Stairway | 482 | 5 | 0 | 2 | 1 |
+-----+-----+-----+-----+-----+-----+-----+
(2 rows)
```

ON DELETE Choices

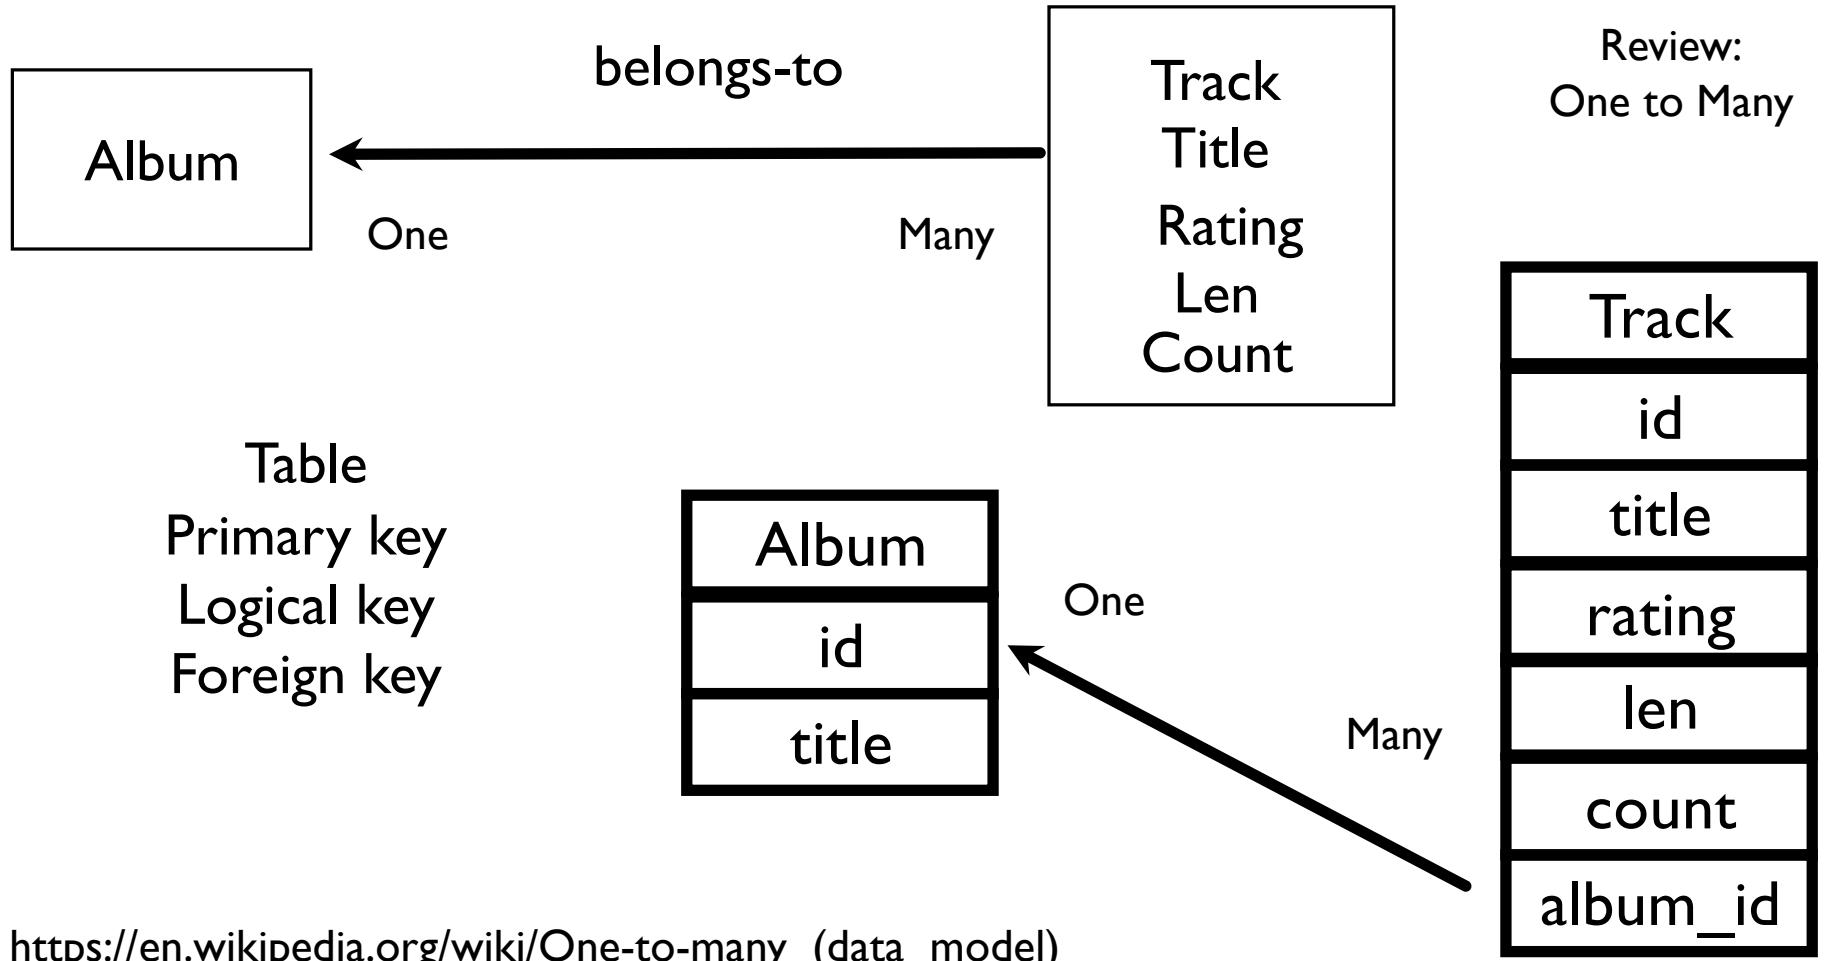
- Default / RESTRICT – Don't allow changes that break the constraint
- CASCADE – Adjust child rows by removing or updating to maintain consistency
- SET NULL – Set the foreign key columns in the child rows to null

<http://stackoverflow.com/questions/1027656/what-is-mysqls-default-on-delete-behavior>

Many-to-Many Relationships



www.tsugi.org



[https://en.wikipedia.org/wiki/One-to-many_\(data_model\)](https://en.wikipedia.org/wiki/One-to-many_(data_model))

```
music=> SELECT * FROM track;
```

id	title	len	rating	count	album_id	genre_id
1	Black Dog	297	5	0	2	1
2	Stairway	482	5	0	2	1
3	About to Rock	313	5	0	1	2
4	Who Made Who	207	5	0	1	2

One

Many

```
music=> SELECT * FROM genre;
```

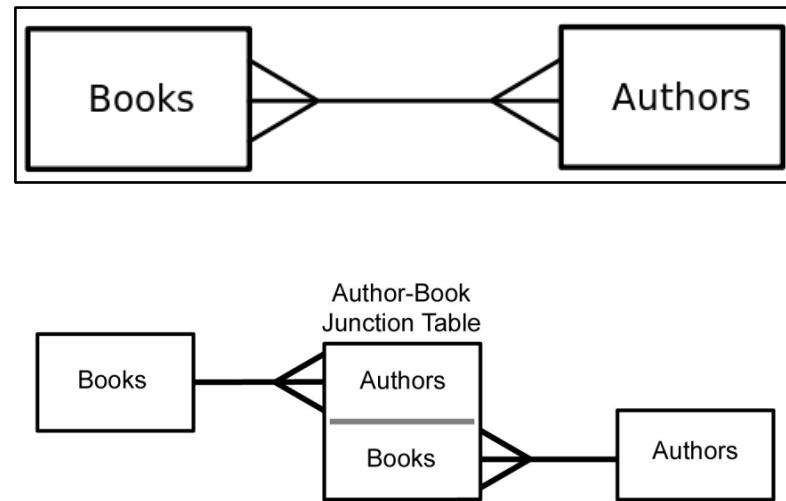
id	name
1	Rock
2	Metal

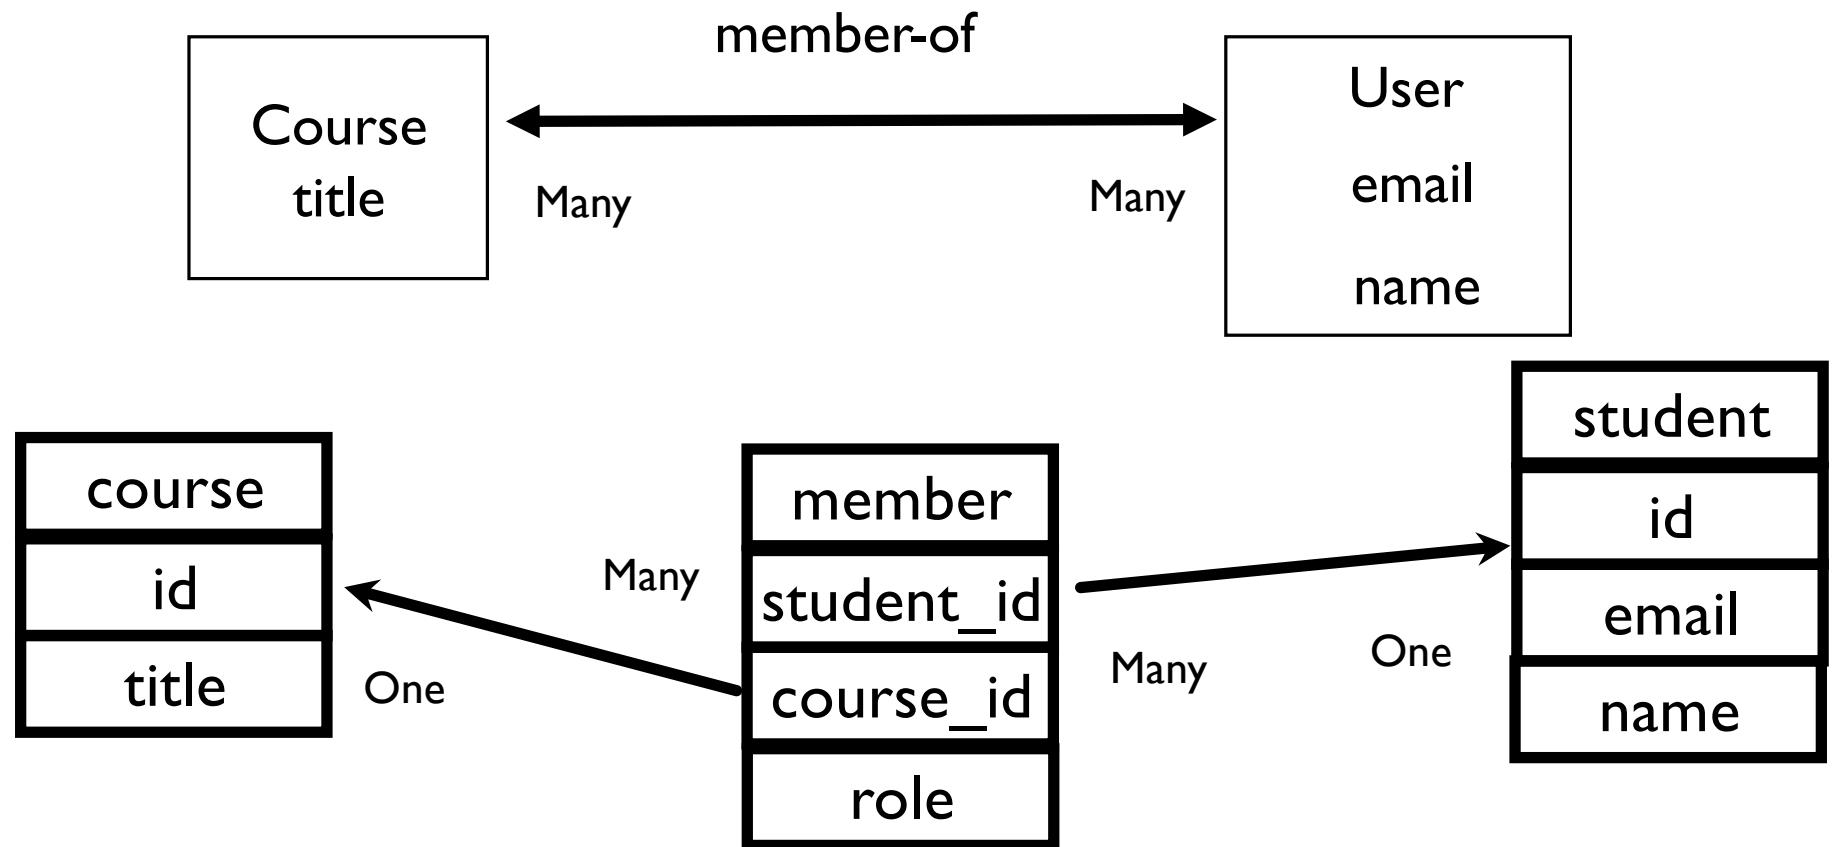


[https://en.wikipedia.org/wiki/One-to-many_\(data_model\)](https://en.wikipedia.org/wiki/One-to-many_(data_model))

Many to Many

- Sometimes we need to model a relationship that is many to many.
- We need to add a “connection” table with two foreign keys.
- There is usually no separate primary key.



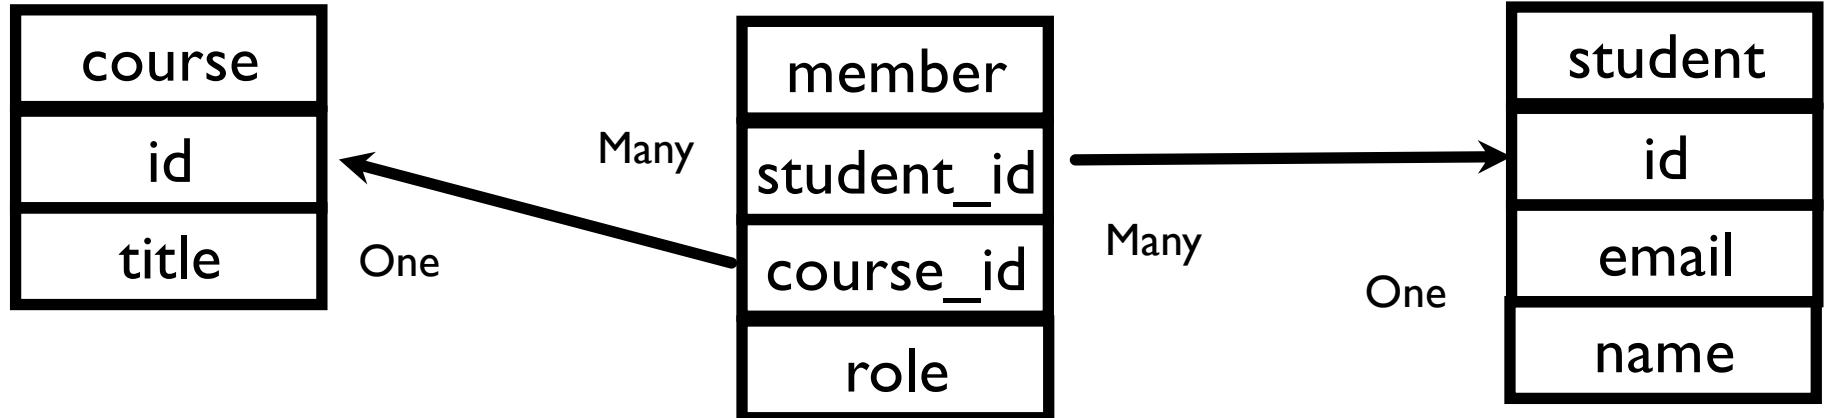


[https://en.wikipedia.org/wiki/Many-to-many_\(data_model\)](https://en.wikipedia.org/wiki/Many-to-many_(data_model))

Start with a Fresh Database

```
CREATE TABLE student (
    id SERIAL,
    name VARCHAR(128),
    email VARCHAR(128) UNIQUE,
    PRIMARY KEY(id)
) ;
```

```
CREATE TABLE course (
    id SERIAL,
    title VARCHAR(128) UNIQUE,
    PRIMARY KEY(id)
) ;
```



```
CREATE TABLE member (
    student_id INTEGER REFERENCES student(id) ON DELETE CASCADE,
    course_id INTEGER REFERENCES course(id) ON DELETE CASCADE,
    role      INTEGER,
    PRIMARY KEY (student_id, course_id)
);
```

Insert Users and Courses

```
music=> INSERT INTO student (name, email) VALUES ('Jane', 'jane@tsugi.org');
music=> INSERT INTO student (name, email) VALUES ('Ed', 'ed@tsugi.org');
music=> INSERT INTO student (name, email) VALUES ('Sue', 'sue@tsugi.org');
music=> SELECT * FROM student;
+----+----+
| id | name | email |
+----+----+
| 1  | Jane | jane@tsugi.org |
| 2  | Ed   | ed@tsugi.org  |
| 3  | Sue  | sue@tsugi.org |
+----+----+



music=> INSERT INTO course (title) VALUES ('Python');
music=> INSERT INTO course (title) VALUES ('SQL');
music=> INSERT INTO course (title) VALUES ('PHP');
music=> SELECT * FROM COURSE;
+----+----+
| id | title |
+----+----+
| 1  | Python |
| 2  | SQL    |
| 3  | PHP    |
+----+----+
```

Insert Memberships

```
music=> SELECT * FROM student;
+-----+-----+
| id | name | email |
+-----+-----+
| 1  | Jane | jane@tsugi.org |
| 2  | Ed   | ed@tsugi.org  |
| 3  | Sue  | sue@tsugi.org |
```

```
music=> SELECT * FROM course;
+-----+-----+
| id | title |
+-----+-----+
| 1  | Python |
| 2  | SQL    |
| 3  | PHP    |
```

```
INSERT INTO member (student_id, course_id, role) VALUES (1, 1, 1);
INSERT INTO member (student_id, course_id, role) VALUES (2, 1, 0);
INSERT INTO member (student_id, course_id, role) VALUES (3, 1, 0);

INSERT INTO member (student_id, course_id, role) VALUES (1, 2, 0);
INSERT INTO member (student_id, course_id, role) VALUES (2, 2, 1);

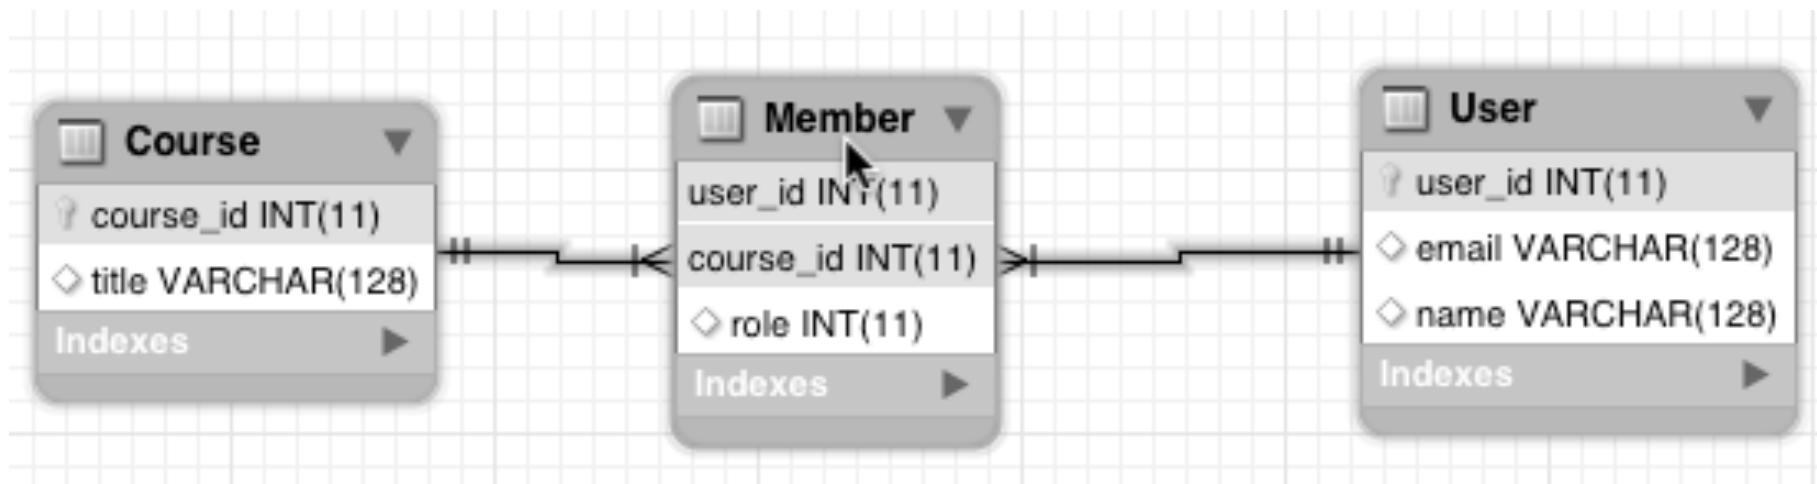
INSERT INTO member (student_id, course_id, role) VALUES (2, 3, 1);
INSERT INTO member (student_id, course_id, role) VALUES (3, 3, 0);
```

```
music=> SELECT * FROM student;
+----+----+-----+
| id | name |      email |
+----+----+-----+
| 1  | Jane | jane@tsugi.org |
| 2  | Ed   | ed@tsugi.org  |
| 3  | Sue  | sue@tsugi.org |
+----+----+-----+
```

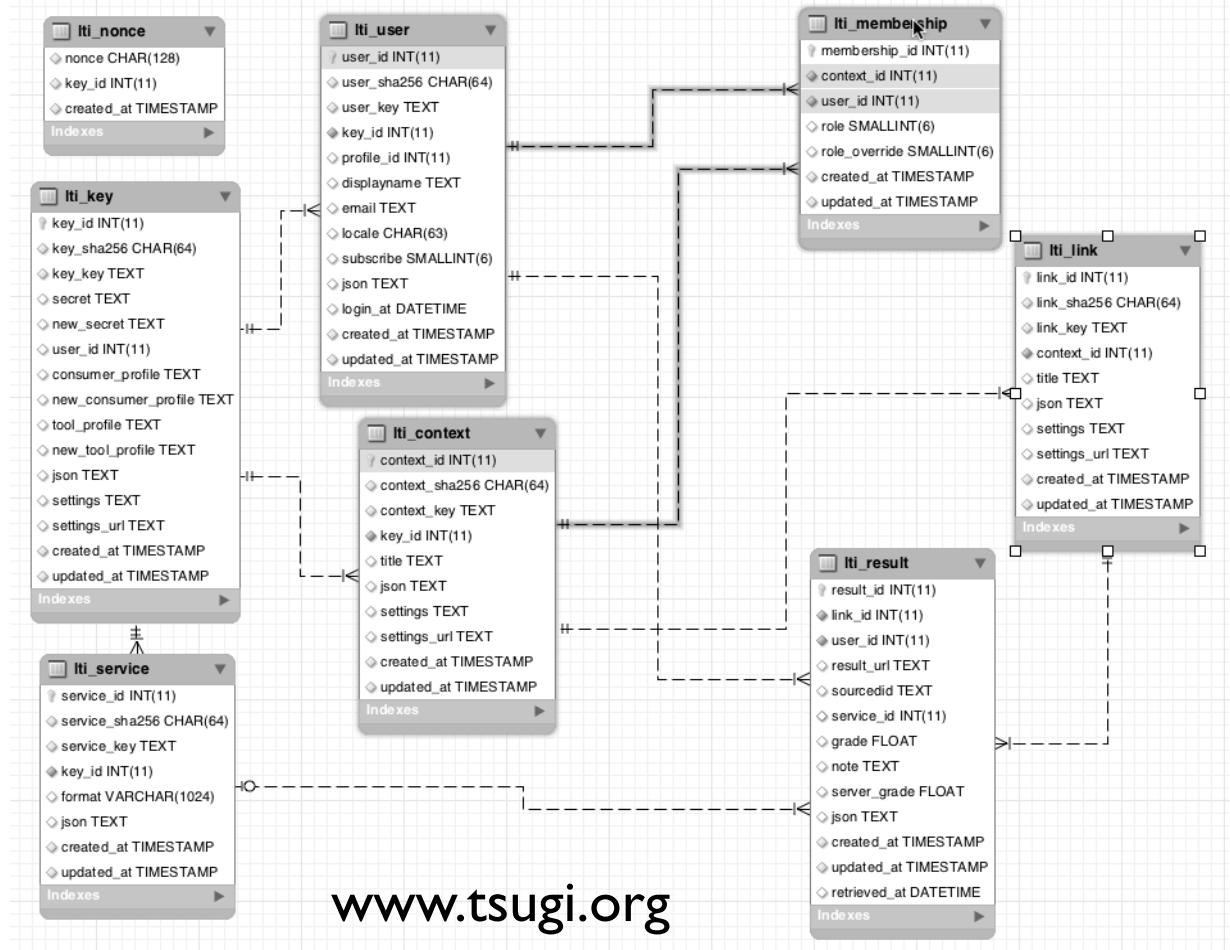
```
music=> SELECT * FROM course;
+----+-----+
| id | title |
+----+-----+
| 1  | Python |
| 2  | SQL    |
| 3  | PHP    |
+----+-----+
```

```
music=> SELECT * FROM member;
+----+----+-----+
| student_id | course_id | role |
+----+----+-----+
| 1           | 1          | 1    |
| 2           | 1          | 0    |
| 3           | 1          | 0    |
| 1           | 2          | 0    |
| 2           | 2          | 1    |
| 2           | 3          | 1    |
| 3           | 3          | 0    |
+----+----+-----+
```

```
music=> SELECT student.name, member.role, course.title
music-> FROM student
music-> JOIN member ON member.student_id = student.id
music-> JOIN course ON member.course_id = course. id
music-> ORDER BY course.title, member.role DESC,
student.name;
      name | role | title
-----+-----+
      Ed   |    1 | PHP
     Sue  |    0 | PHP
    Jane |    1 | Python
      Ed  |    0 | Python
     Sue |    0 | Python
      Ed  |    1 | SQL
    Jane |    0 | SQL
( 7 rows )
```



<https://www.mysql.com/products/workbench/>



www.tsugi.org

Complexity Enables Speed

- Complexity makes speed possible and allows you to get very fast results as the data size grows.
- By normalizing the data and linking it with integer keys, the overall amount of data which the relational database must *scan* is far lower than if the data were simply flattened out.
- It might seem like a tradeoff - spend some time designing your database so it continues to be fast when your application is a success.

Summary

- Relational databases allow us to scale to very large amounts of data.
- The key is to have one copy of any data element and use relations and joins to link the data to multiple places.
- This greatly reduces the amount of data that must be scanned when doing complex operations across large amounts of data.
- Database and SQL design is a bit of an art form.

Acknowledgements / Contributions



These slides are Copyright 2010- Charles R. Severance (www.dr-chuck.com) as part of www.pg4e.com and made available under a Creative Commons Attribution 4.0 License. Please maintain this last slide in all copies of the document to comply with the attribution requirements of the license. If you make a change, feel free to add your name and organization to the list of contributors on this page as you republish the materials.

Continue new Contributors and Translators here

Initial Development: Charles Severance, University of Michigan
School of Information

Insert new Contributors and Translators here including names and dates